

World's dumbest clinical data service

- Healthy/Sick
- Authorizations needed stored alongside data, in another table
- Very simple SQL-as-service: resource (Individuals), select, where
- Separate Consents/Authorizations table: allows single individual to be consented to/part of multiple projects (MoHN)



Individuals

id	status
P001	Healthy
P002	Sick
P003	Healthy
P004	Sick
P005	Healthy
P006	Sick


Consents

id	project	consent
P001	profyle	TRUE
P002	profyle	TRUE
P003	profyle	TRUE
P004	tf4cn	TRUE
P005	tf4cn	TRUE
P006	tf4cn	TRUE

World's dumbest clinical data service

- Higher level services call the sqlite3 service
 - Clinical service
 - /individuals
 - /individuals/P001
 - Analytics service
 - /n
 - /healthy_fraction

clinical_service



`select individuals.*`

analytics_service



`select count(id)
...
where status = 'Healthy'`



Individuals


id	status
P001	Healthy
P002	Sick
P003	Healthy
P004	Sick
P005	Healthy
P006	Sick

Consents

id	project	consent
P001	profyle	TRUE
P002	profyle	TRUE
P003	profyle	TRUE
P004	tf4cn	TRUE
P005	tf4cn	TRUE
P006	tf4cn	TRUE

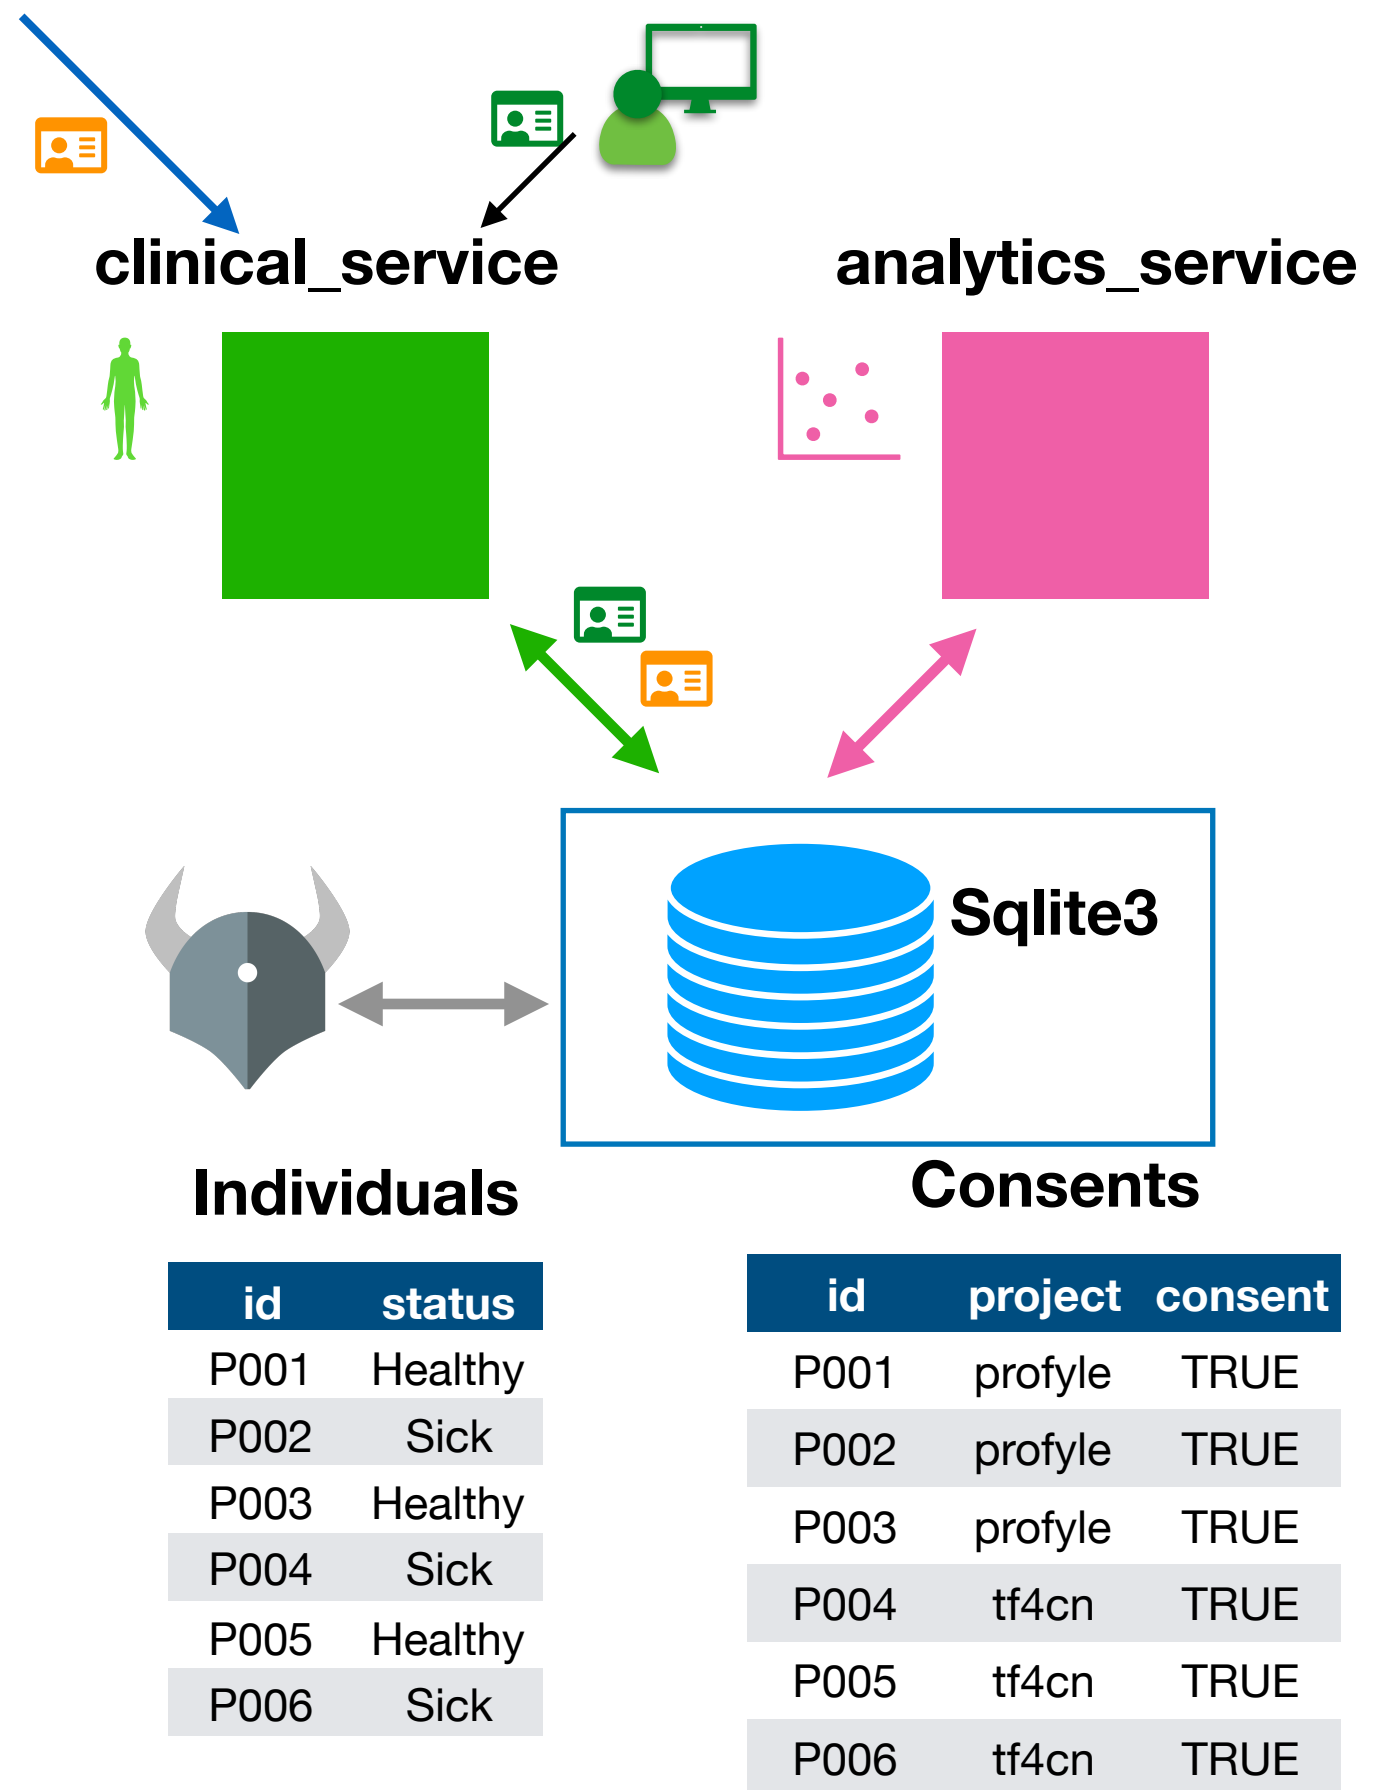
How is request made?

- Request comes in with ID token
- Tyk @ data host requests DAC claims tokens based on that identity from remote DAC portals if they exist
- Not all datasets will have a DAC portal - also support authorization lists




Authorization via OPA

- Tokens (ID token, DAC claims tokens) are passed through the services
- Closest to data, data service queries OPA for authorization info




Authorization via OPA

- Sends query to OPA service
- Service evaluates based on policy, and sends yes/no/“under these conditions” back



```
input = {  
  user: 🧑,  
  entitlements: [ {'profile_member', 🧑} ],  
  method: "GET",  
  path: ["individuals"],  
  unknowns: ["individuals", "consents"]  
}
```



```
allow = True  
or  
allow = False  
or  
allow if....
```

If any of the tokens are expired, or not signed by right key/modified, authorization fails

```
# get the user (subject) from the identity token
idtoken := {"payload": payload} { io.jwt.decode(input.user, [_, payload, _]) }
subject := idtoken.payload.sub
```

```
default valid_tokens = false
default valid_id_token = false
default valid_entitlement_tokens = false
default allow = false
```

If any of the tokens are expired, or not signed by right key/modified, authorization fails

```
# authorization is denied if the any token is expired or if signature fails validation
# or if entitlement (DAC claim) token doesn't match ID token
# this is also where we'd check that the issuer is one of
# our trusted federation partners, etc.
valid_id_token = true {
  [id_valid, id_header, id_payload] := io.jwt.decode_verify(input.user, {"secret": "secret"})
  id_valid == true
}

# if there are no entitlement tokens, they are valid (there's no invalid ones)
valid_entitlement_tokens = true {
  count(input.entitlements) == 0
}

valid_entitlement_tokens = true {
  # if any of the claims tokens are expired or fail signature validation
  # or the subject isn't the same as the ID subject
  # then we fail
  entitlement_checks := [{"valid":valid, "payload":payload} |
    [valid, header, payload] := io.jwt.decode_verify(input.entitlements[_].jwt, {"secret": "secret"})]
  all([entitlement_checks[_].valid == true])
  all([entitlement_checks[_].payload.sub == subject])
}

valid_tokens {
  valid_id_token
  valid_entitlement_tokens
}
```

**If passes valid_tokens check,
and requesting a particular item,
succeed if entitlements include consents
for that item**

```
## access list for projects that don't have a DAC
access_list = {"alice":["tf4cn_member"],
               "bob":  ["tf4cn_member"]}
```

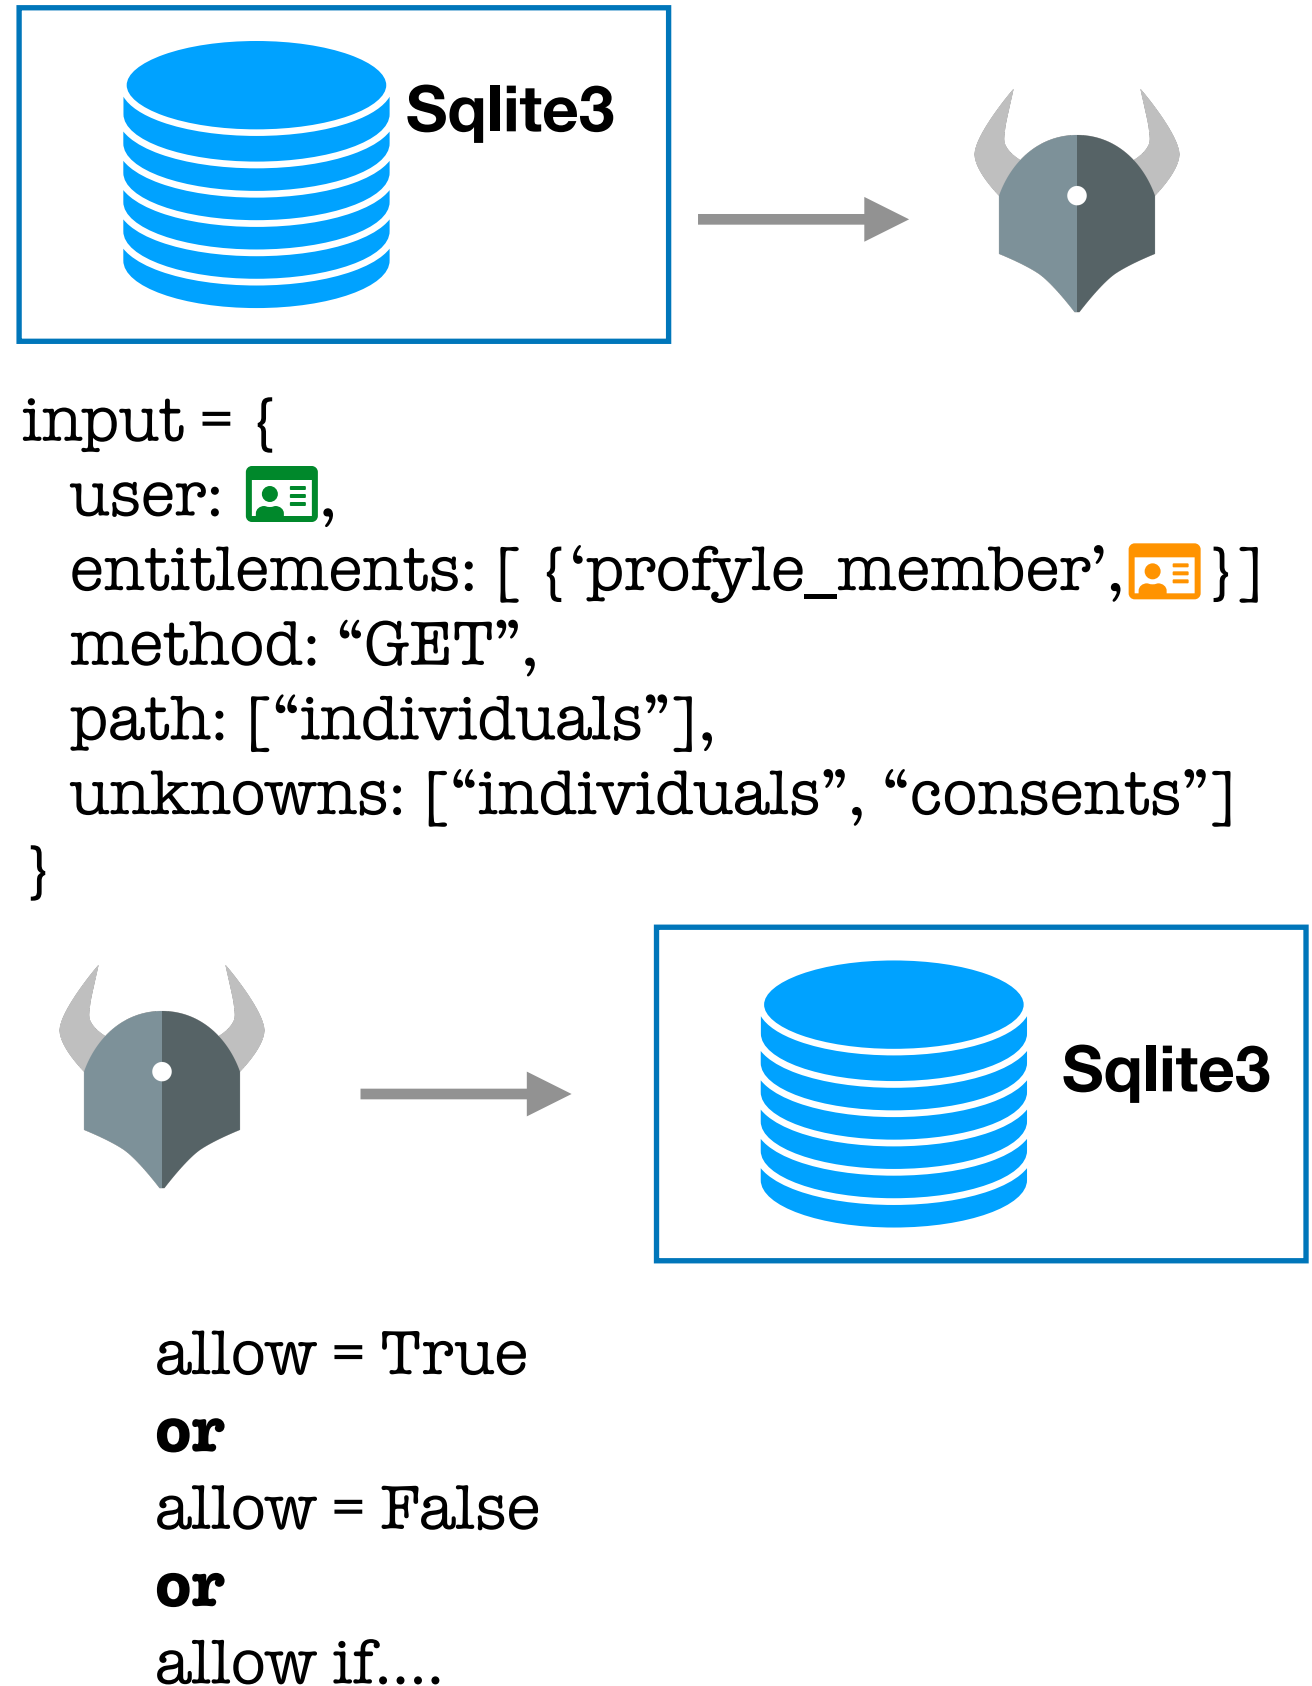
```
# get set of all entitlements from tokens and access list
claim_entitlements := { name | name = input.entitlements[i].name;
                       [ct_head, ct_payload, ct_sig] = io.jwt.decode(input.entitlements[i].jwt);
                       ct_payload[name] == true }
access_list_entitlements := { item | access_list[subject][_ ] = item }
entitlements := union({claim_entitlements, access_list_entitlements})
```


**If passes valid_tokens check,
and requesting a particular item,
succeed if entitlements include consents
for that item**

```
# authorize a single item
allow = true {
  input.method = "GET"
  input.path = ["individuals", iid]
  # for a single item we can just be given the consents and make the authorization decision
  any({input.consents[i] == entitlements[j]})
}
```


Authorization via OPA

- For valid tokens and a single item, know all relevant information at query time
- Expect a True/False back




Authorization via OPA

- For a listing of items, don't know what items are yet.
- Could:
 - get all items
 - check individually
- Or..



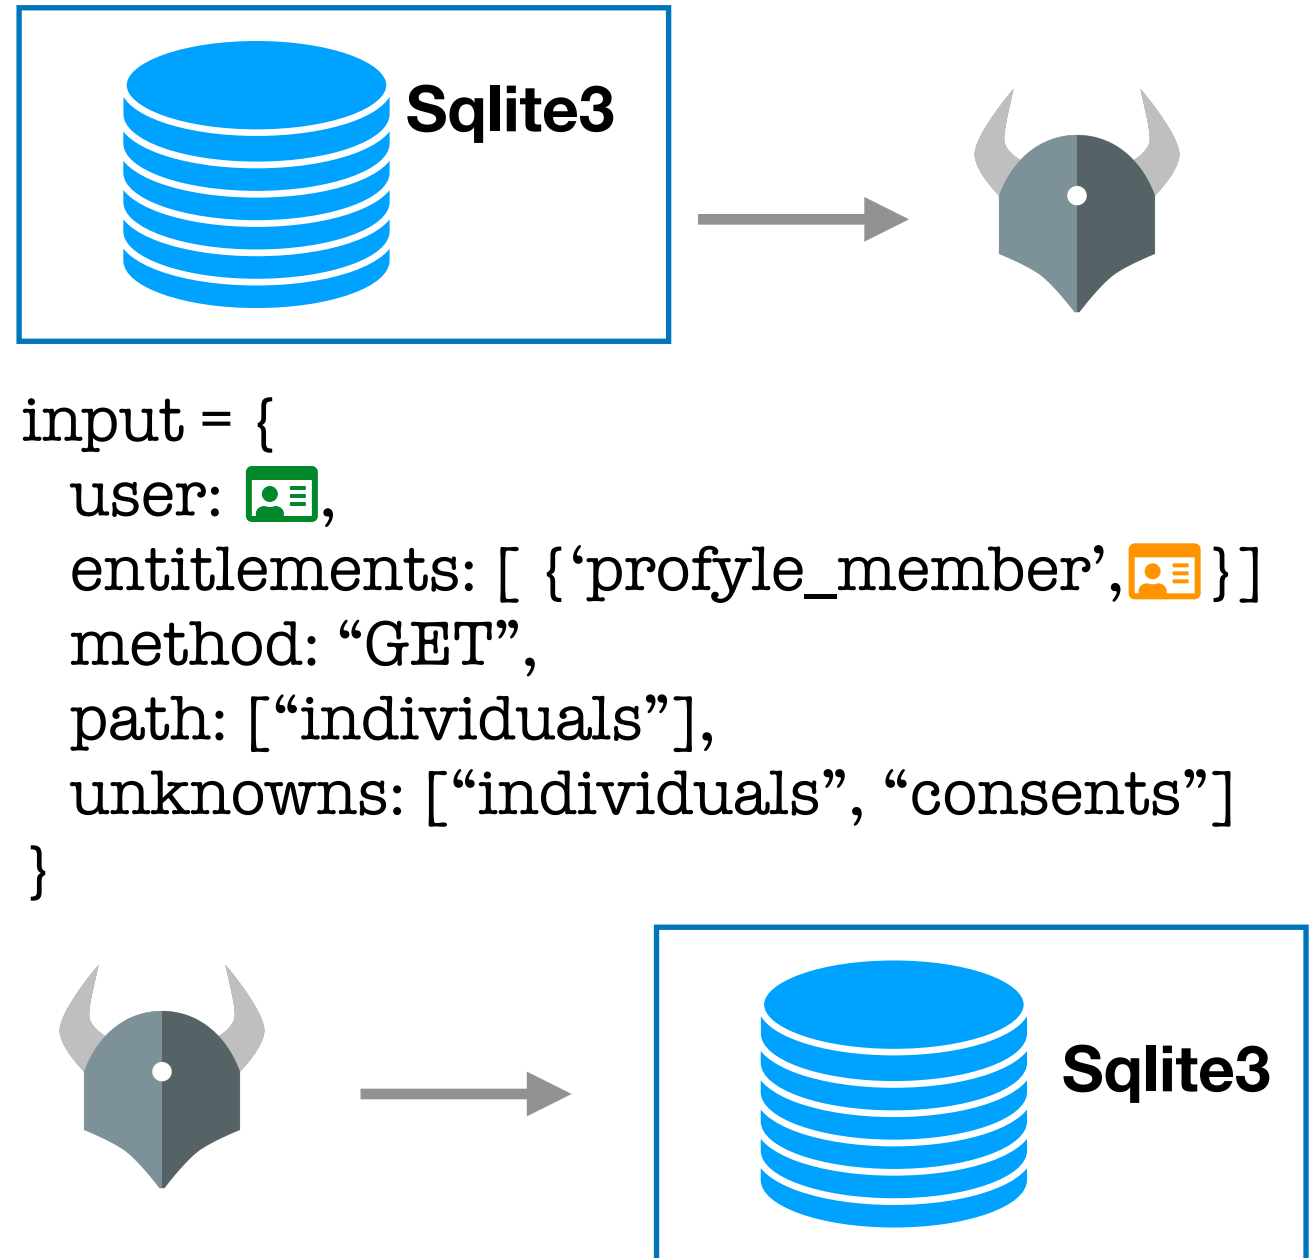
```
input = {  
  user: 🧑,  
  entitlements: [ {'profile_member', 🧑} ],  
  method: "GET",  
  path: ["individuals"],  
  unknowns: ["individuals", "consents"]  
}
```



```
allow = True  
or  
allow = False  
or  
allow if....
```

Authorization via OPA

- If the response is a “yes, if”, is sent as a series of conditions
- Server translates this into SQL clauses
- Clauses are sanitized (by SQLAlchemy) and inserted into query



A row is allowed if an individual has a consent that matches the DAC entitlements,

```
# authorize items from a list
allow = true {
  input.method = "GET"
  input.path = ["individuals"]

  some x
  row_allowed[x]
}

row_allowed[x] = true {
  data.individuals[x].id = data.consents[x].id
  data.consents[x].consent = true
  proj := data.consents[x].project

  input.entitlements[i].name == proj
  [_, payload, _] := io.jwt.decode(input.entitlements[i].jwt)
  payload[proj] == true
}
```

...or, row is also allowed
if an individual has a consent that
matches locally-declared entitlements.

```
## access list for projects that don't have a DAC
access_list = {"alice": ["tf4cn_member"],
               "bob":  ["tf4cn_member"]}
```


```
# Alternately, item is allowed if the data consent matches an entitlement in access list above
row_allowed[x] = true {
  data.individuals[x].id = data.consents[x].id
  data.consents[x].consent = true
  some j
  data.consents[x].project = access_list[subject][j]
}
```

(Have to be a bit more explicit
for current SQL translation to work)


Authorization via OPA

- Higher level services call the sqlite3 service
- Don't have to make any authz checks
 - But probably good idea to check token existence, validity
- Authz enforced by sqlite service following OPA-defined policy
- Unauthorized rows never leave database

clinical_service



analytics_service



Individuals

id	status
P001	Healthy
P002	Sick
P003	Healthy
P004	Sick
P005	Healthy
P006	Sick

Consents

id	project	consent
P001	profyle	TRUE
P002	profyle	TRUE
P003	profyle	TRUE
P004	tf4cn	TRUE
P005	tf4cn	TRUE
P006	tf4cn	TRUE